

## Dog Breed Classifier

### Table of content

<b>Table of content</b>	<b>1</b>
<b>Definition</b>	<b>2</b>
Project Overview	2
<b>Problem Statement</b>	<b>2</b>
Metrics	2
<b>Analysis</b>	<b>3</b>
Data Exploration and Visualization	3
Human dataset	3
Sample images with labels	4
Dog dataset	6
Sample images with labels	7
Algorithms and Techniques	7
Benchmark	8
<b>Methodology</b>	<b>8</b>
Data Pre-processing	8
Implementation	8
The classifier training stage	8
The network architecture	9
Using the trained model to classify dogs	10
<b>Results</b>	<b>11</b>
Model Evaluation and Validation	11
<b>Conclusion</b>	<b>12</b>
Reflection	12
Improvement	13

# Definition

## Project Overview

A dog breed classifier is a computer vision classifier that uses Convolutional Neural Networks (CNN) to build a pipeline to process real-world, user-supplied images.

Given an image of a dog, The algorithm will identify an estimate of the canine's breed. If supplied an image of a human, the code will identify the resembling dog breed.

## Problem Statement

The goal is to create a dog breed classifier based on a dog or human image.

In the real world, there are a lot of different dog breeds. While Humans can detect few dog breeds, The classifier can detect more than 100 different dog breeds.

The output will be the Input Image labeled with the predicted dog class.  
The Input will be an RGB image of a dog or a human.

Machine learning will be used to capture the features for the Input image and use those features to predict the dog class.

The project will be developed by completing the following tasks:

1. Download the input data
2. Preprocess the data and see If the classes are imbalanced
3. Create Human Face detection using open cv library haar cascades
4. Create a general dog detector using a pre-trained model like VGG which is a sequence of convolution and pooling layers
5. Using transfer learning to create the dog breed classifier after dropping the classifier layer in the pre-trained model and attach the new one after changing the output number.
6. Take a raw image and return the dog breed.

## Metrics

### CrossEntropyLoss

It is useful when training a classification problem with C classes. This Cross-Entropy Loss is useful when you have an unbalanced training set.

$$\text{loss}(x, \text{class}) = -\log \left( \frac{\exp(x[\text{class}])}{\sum_j \exp(x[j])} \right) = -x[\text{class}] + \log \left( \sum_j \exp(x[j]) \right)$$

The losses are averaged across observations for each minibatch.

$$\text{loss} = \frac{\sum_{i=1}^N \text{loss}(i, \text{class}[i])}{\sum_{i=1}^N \text{weight}[\text{class}[i]]}$$

### Accuracy

is a common metric for binary and multi-class classifiers; it takes into account both true positives and true negatives with equal weight.

accuracy =(true positives + true negative) / size of dataset

Because we have multi-class classification precision and recall will not be a good metric as we are focus on the total number of correct prediction

We can use ROC-AUC as a metric and compare the results as the next Improvements.

# Analysis


## Data Exploration and Visualization

### Human dataset

The human dataset has 13233 images with 5749 unique Humans. George\_W\_Bush has the most counts with 530 image

The top 20 Humans were:

Name	# of images	George.W.Bush
George_W_Bush	530	
Colin_Powell	236	
Tony_Blair	144	
Donald_Rumsfeld	121	
Gerhard_Schroeder	109	
Ariel_Sharon	77	
Hugo_Chavez	71	
Junichiro_Koizumi	60	
Jean_Chretien	55	
John_Ashcroft	53	
Serena_Williams	52	
Jacques_Chirac	52	
Vladimir_Putin	49	
Luiz_Inacio_Lula_da_Silva	48	
Gloria_Macapagal_Arroyo	44	
Jennifer_Capriati	42	
Arnold_Schwarzenegger	42	
Laura_Bush	41	
Lleyton_Hewitt	41	
Hans_Blix	39	




## Sample images with labels



Nabil\_Shaath



Brendan\_Stai



Jennifer\_Garner




Lisa\_Murkowski



Jamie\_Cooke




Hernan\_Crespo



Cemil\_Cicek



Allison\_Janney



Travis\_Rudolph



Demetrius\_Ferraciu




Sanjay\_Gupta




Brian\_Cook



Hussam\_Mohammed\_Amin



LeRoy\_Millette\_Jr



Eddie\_Lucio



Ward\_Cuff




Nia\_Vardalos



Don\_King



Bill\_Bradley





Steve\_Zahn

## Dog dataset

The Dog dataset has 8351 images with 133 unique dog breeds. the Alaskan malamute has the most counts with 95 image

The top 20 Dogs were:


Dog breed	# counts	Alaskan malamute
Alaskan_malamute	96	
Border_collie	93	
Basset_hound	92	
Dalmatian	89	
Bull_terrier	87	
Bullmastiff	86	
Basenji	86	
Cavalier_king_charles_spaniel	84	
Australian_cattle_dog	83	
Australian_shepherd	83	
Dachshund	82	
Irish_terrier	82	
American_staffordshire_terrier	82	
Boston_terrier	81	
Briard	81	
Bernese_mountain_dog	81	
Affenpinscher	80	
American_eskimo_dog	80	
Bloodhound	80	
Cane_corso	80	



## Sample images with labels



Lowchen




Cocker\_spaniel




Cairn\_terrier



Kuvasz



Kerry\_blue\_terrier



Pointer



Norfolk\_terrier




American\_staffordshire\_terrier



Boykin\_spaniel



Welsh\_springer\_spaniel



Mastiff



Yorkshire\_terrier



Alaskan\_malamute




Irish\_water\_spaniel




Dandie\_dinmont\_terrier



Briard




Leonberger



Italian\_greyhound



Neapolitan\_mastiff



Irish\_wolfhound

## Algorithms and Techniques

The classifier is a Convolutional Neural Network, which is the state-of-the-art algorithm for most image processing tasks, including classification. It needs a large amount of training data compared to other approaches.

The following parameters can be tuned to optimize the classifier:

1. Training parameters:
  - Training length (number of epochs)
  - Batch size (how many images to look at once during a single training step)
  - Learning rate (how fast to learn; this can be dynamic)
2. Neural network architecture:
  - Number of layers



- Layer types ( convolutional, fully-connected, or pooling )
3. Preprocessing parameters

During training, both the training and the validation sets are loaded into the RAM. After that, random batches are selected to be loaded into the GPU memory for processing.

We could use Two Techniques:

1- Train the model from scratch :

We can build our model from scratch specifically to the problem. This Technique is when we want to build a detector for simple use cases

The problem with that Technique that we need to train the model from the beginning and need to design every layer in the model (Convolution, fully connected and pooling ...etc)

2- Use pre-trained Network by transfer learning

We can use the CNN model that has been trained and use the layers that detect the features from the input whatever the Input type or shape and replace the pre-trained Network classifier for a custom classifier.

This Technique save a lot of training time as we freeze the feature layers and train only the classifier

I have used transfer learning in the final solution using the VGG16 model as It was trained on a huge dataset.

## Benchmark

I have used Two benchmarks.

the first one was training a model from scratch and test the model and record the highest score that we can get with a given dataset and try to beat this result using other techniques.

The second one was to set an accuracy threshold from the model that we can accept for the application. I targeted 75% accuracy on the test set

I have set a timing constraint for real-time classification, so the model target to classify the input image in time doesn't exceed 50 ms



# Methodology

## Data Pre-processing

The preprocessing has the following steps:

1. Make augmentation for images (RandomRotation, RandomResizedCrop and RandomHorizontalFlip)
2. Normalize the input images
3. The images are divided into training, validation, and test sets
4. Create the data loaders by configuring the batch size and shuffle the training data

## Implementation

The implementation process can be split into two main stages:

1. The classifier training stage
2. Using the trained model to classify dogs

### The classifier training stage

During the first stage, the classifier was trained on the preprocessed training data using transfer learning.

The training steps were:

1. load both the training and validation images into memory using the data loaders that we have created
2. Define the network architecture and training parameters
3. Define the loss function, accuracy
4. Train the network, logging the validation/training loss and the validation accuracy
5. If the accuracy is not high enough, return to step 2
6. Save and freeze the trained network

### The network architecture

I have used transfer learning using the vgg16 model

```
(features): Sequential(
  (0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (1): ReLU(inplace=True)
  (2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (3): ReLU(inplace=True)
  (4): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
```

```

(5): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(6): ReLU(inplace=True)
(7): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(8): ReLU(inplace=True)
(9): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
(10): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(11): ReLU(inplace=True)
(12): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(13): ReLU(inplace=True)
(14): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(15): ReLU(inplace=True)
(16): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
(17): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(18): ReLU(inplace=True)
(19): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(20): ReLU(inplace=True)
(21): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(22): ReLU(inplace=True)
(23): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
(24): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(25): ReLU(inplace=True)
(26): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(27): ReLU(inplace=True)
(28): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(29): ReLU(inplace=True)
(30): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
)
(avgpool): AdaptiveAvgPool2d(output_size=(7, 7))
(classifier): Sequential(
  (fc1): Linear(in_features=25088, out_features=4096, bias=True)
  (relu): ReLU()
  (Dropout1): Dropout(p=0.5, inplace=False)
  (fc2): Linear(in_features=4096, out_features=1024, bias=True)
  (Dropout2): Dropout(p=0.5, inplace=False)
  (fc3): Linear(in_features=1024, out_features=133, bias=True)
  (output): LogSoftmax()
)
)

```

I have kept the features layers as It has already trained to detect shapes and patterns and replaced the classifiers layers with one that I need to classify the new 133 dog breed.

The new classifier becomes:

```

(classifier): Sequential(
  (fc1): Linear(in_features=25088, out_features=4096, bias=True)
  (relu): ReLU()
  (Dropout1): Dropout(p=0.5, inplace=False)

```

```
(fc2): Linear(in_features=4096, out_features=1024, bias=True)
(Dropout2): Dropout(p=0.5, inplace=False)
(fc3): Linear(in_features=1024, out_features=133, bias=True)
(output): LogSoftmax()
)
```

## Using the trained model to classify dogs

1. Create the model using the given layers
2. Load the pre-trained weights
3. Process the input Images
4. Using open cv to detect faces in Images
5. Using a pre-trained VGG model to detect dogs
6. Predict the input Image using the model

# Results

## Model Evaluation and Validation

During development, a validation set was used to evaluate the model.

The final architecture and hyperparameters were chosen because they performed the best among the tried combinations.

For a complete description of the final model and the training process, refer to The network architecture section

I have trained the model after replacing the Classifier layers for 10 Epochs with the following train and Validation losses:

Epoch: 1	Training Loss: 1.343471	Validation Loss: 0.829716
Epoch: 2	Training Loss: 1.270921	Validation Loss: 0.739991
Epoch: 3	Training Loss: 1.208083	Validation Loss: 0.681313
Epoch: 4	Training Loss: 1.185337	Validation Loss: 0.700088

Epoch: 5	Training Loss: 1.175069	Validation Loss: 0.613085
Epoch: 6	Training Loss: 1.112700	Validation Loss: 0.649210
Epoch: 7	Training Loss: 1.109303	Validation Loss: 0.549573
Epoch: 8	Training Loss: 1.084382	Validation Loss: 0.585246
Epoch: 9	Training Loss: 1.055470	Validation Loss: 0.690492
Epoch: 10	Training Loss: 1.021630	Validation Loss: 0.687842

So as we can see the Validation was still decreasing so we can have better accuracy If we Increased the number of epochs but It will take a longer time.

Using the test set, the Test Accuracy was 80% (677/836)

## Justification

The Test accuracy using **transfer learning** is **80%**

- For the first benchmark, The Test accuracy using the trained model from **scratch** was **18%** using the same input data set and environment which is a significant improvement.
- For the second benchmark, the target was to exceed the 75% test accuracy and that was done using **transfer learning**.
- For the timing constraints for real-time, we were target to classify images at average 40ms/image  
It took 5.42 Second to classify 200 Image so the average was 27 ms/image

I used my local machine (6 GByte GPU"  
NVIDIA Corporation TU116M [GeForce GTX 1660 Ti Mobile]", 12 logical CPU  
"Intel® Core™  
i7-9750H CPU @ 2.60GHz × 12") in the training and testing

The result is stronger than I expected as the model was trained for just 10 Epochs and we got 80 % accuracy.


We didn't train the feature layers and use them as it with trained weights so we achieved our target with min. Training time.

The model is good enough to detect 8 images from the given 10 input images. this ratio could be increased by train the model for more Epochs

## Conclusion

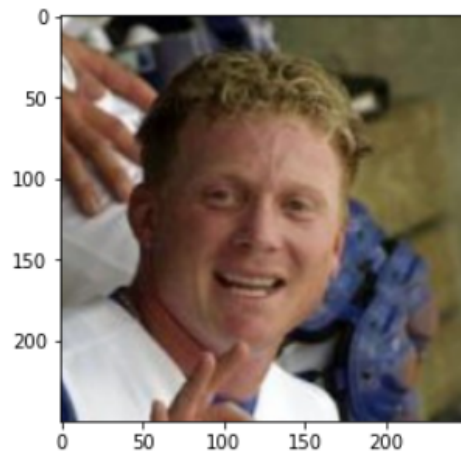
Using Convolution neural networks we can classify the dog breeds with an accuracy of more than 80 % and we can even make this classification for humans.

Hello, Human!




You look like a Dogue de bordeaux


Hello, Human!



You look like a Dachshund



You look like a Affenpinscher



You look like a Brussels griffon

## Reflection

The process used for this project can be summarized using the following steps:

1. An initial problem and relevant, public datasets were found
2. The data was downloaded and preprocessed (segmented)
3. A benchmark was created for the classifier
4. The classifier was trained using the data and transfer learning (multiple times, until a good set of parameters, were found)
5. The model was used to detect the dog breed on unseen images.

## Improvement

To achieve better accuracy for the model

- 1- we can use different model features layers for our model
- 2- we can CNN to detect humans better than a face classifier
- 3- we can train for more epochs